

Enterprise Knowledge Base - Complex RAG Test Document

Section 1: Topic 1

This section describes enterprise AI systems, retrieval augmented generation, vector databases, document ingestion, chunking strategies, metadata filtering, embeddings, reranking, governance, compliance, and deployment. Each paragraph contains repeated but varied information to test semantic retrieval. This section describes enterprise AI systems, retrieval augmented generation, vector databases, document ingestion, chunking strategies, metadata filtering, embeddings, reranking, governance, compliance, and deployment. Each paragraph contains repeated but varied information to test semantic retrieval.

Section 2: Topic 2

This section describes enterprise AI systems, retrieval augmented generation, vector databases, document ingestion, chunking strategies, metadata filtering, embeddings, reranking, governance, compliance, and deployment. Each paragraph contains repeated but varied information to test semantic retrieval. This section describes enterprise AI systems, retrieval augmented generation, vector databases, document ingestion, chunking strategies, metadata filtering, embeddings, reranking, governance, compliance, and deployment. Each paragraph contains repeated but varied information to test semantic retrieval.

Section 3: Topic 3

This section describes enterprise AI systems, retrieval augmented generation, vector databases, document ingestion, chunking strategies, metadata filtering, embeddings, reranking, governance, compliance, and deployment. Each paragraph contains repeated but varied information to test semantic retrieval. This section describes enterprise AI systems, retrieval augmented generation, vector databases, document ingestion, chunking strategies, metadata filtering, embeddings, reranking, governance, compliance, and deployment. Each paragraph contains repeated but varied information to test semantic retrieval.

Section 4: Topic 4

This section describes enterprise AI systems, retrieval augmented generation, vector databases, document ingestion, chunking strategies, metadata filtering, embeddings, reranking, governance, compliance, and deployment. Each paragraph contains repeated but varied information to test semantic retrieval. This section describes enterprise AI systems, retrieval augmented generation, vector databases, document ingestion, chunking strategies, metadata filtering, embeddings, reranking, governance, compliance, and deployment. Each paragraph contains repeated but varied information to test semantic retrieval.

Section 5: Topic 5

This section describes enterprise AI systems, retrieval augmented generation, vector databases, document ingestion, chunking strategies, metadata filtering, embeddings, reranking, governance, compliance, and deployment. Each paragraph contains repeated but varied information to test semantic retrieval. This section describes enterprise AI systems, retrieval augmented generation, vector databases, document ingestion, chunking strategies, metadata filtering, embeddings, reranking, governance, compliance, and deployment. Each paragraph contains repeated but varied information to test semantic retrieval.

Project Inventory

Project	Owner	Vector DB	Embedding	Status
RAG-101	Alice	Weaviate	text-embedding-3-small	Production
DOC-202	Bob	Pinecone	text-embedding-3-large	Testing
CHAT-303	Carol	ChromaDB	bge-base-en	Development
SEARCH-404	David	Milvus	e5-large-v2	Production

Support Tickets

ID	Priority	Issue	Resolution
INC1001	High	Embedding generation failed	Restart embedding service
INC1002	Medium	Chunk overlap incorrect	Update chunking config
INC1003	Critical	Vector DB unavailable	Restore cluster and retry ingestion
INC1004	Low	UI alignment	CSS update

Questions this document can answer:

- Which project uses Weaviate?
- Which embedding model is used by DOC-202?
- What issue affected INC1003?
- Who owns CHAT-303?
- Which projects are in production?